

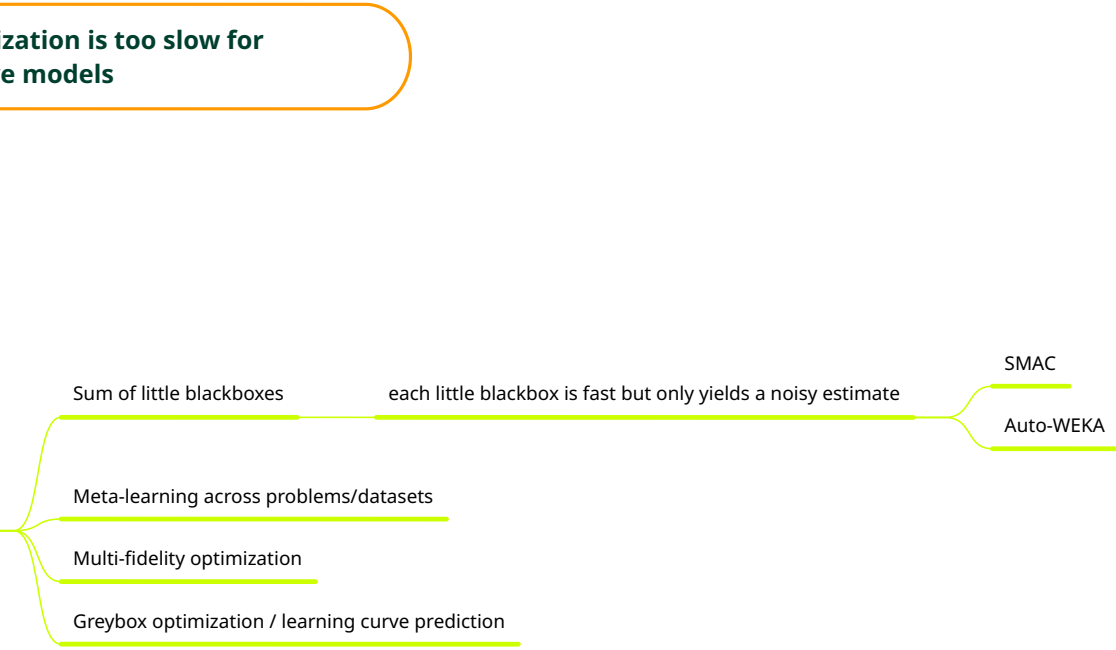
HPO - Speedup

Multi-Fidelity Optimization

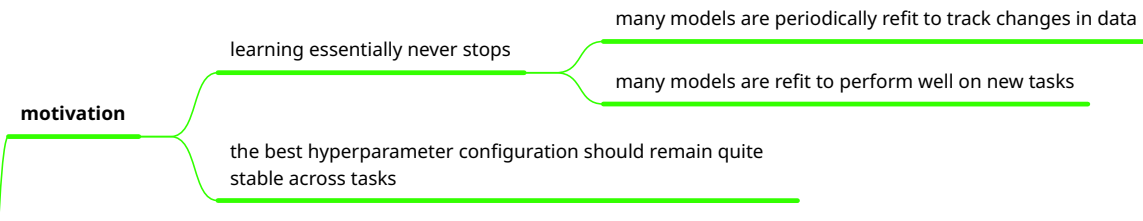


Blackbox optimization is too slow for tuning expensive models

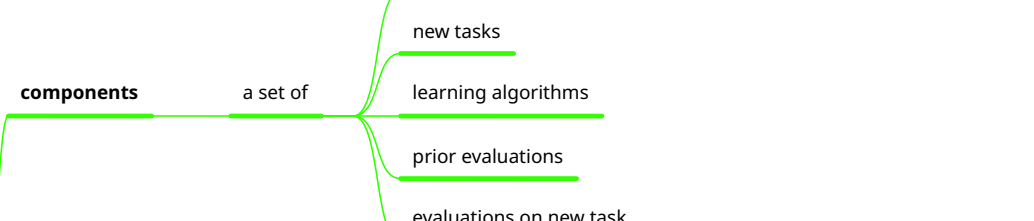
Methods




motivation



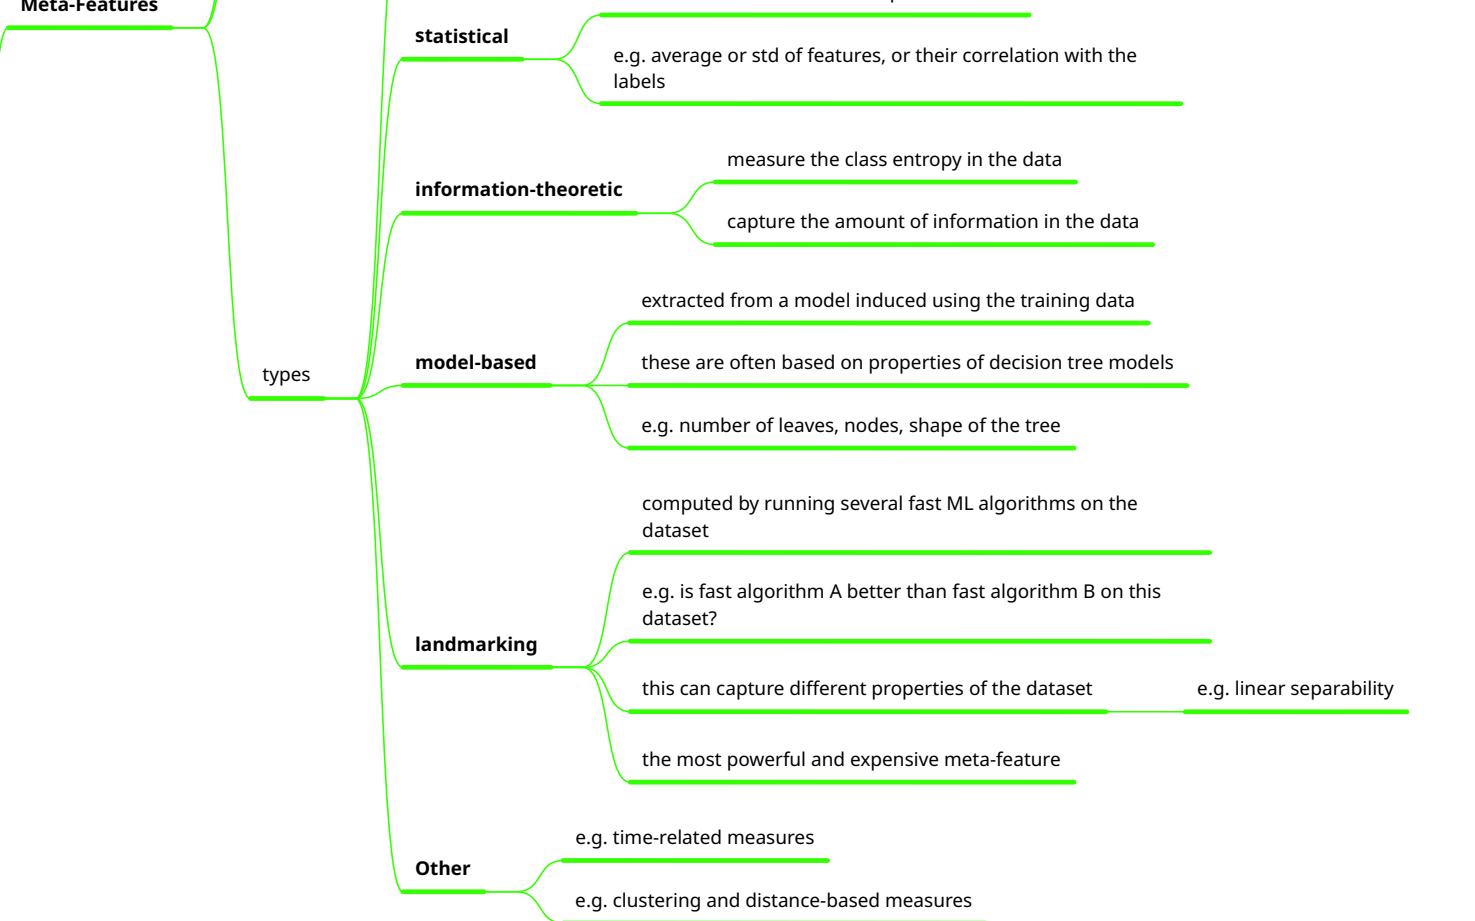
components



goal

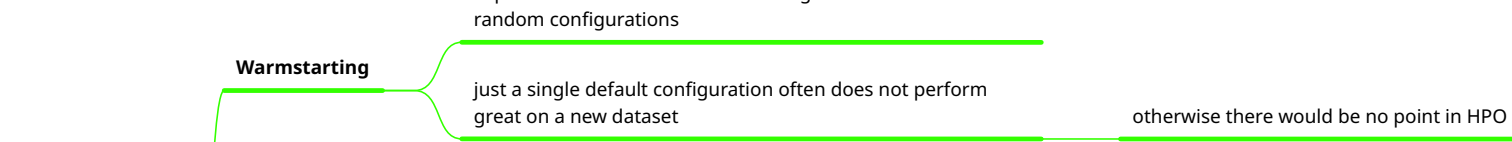


Meta-Features

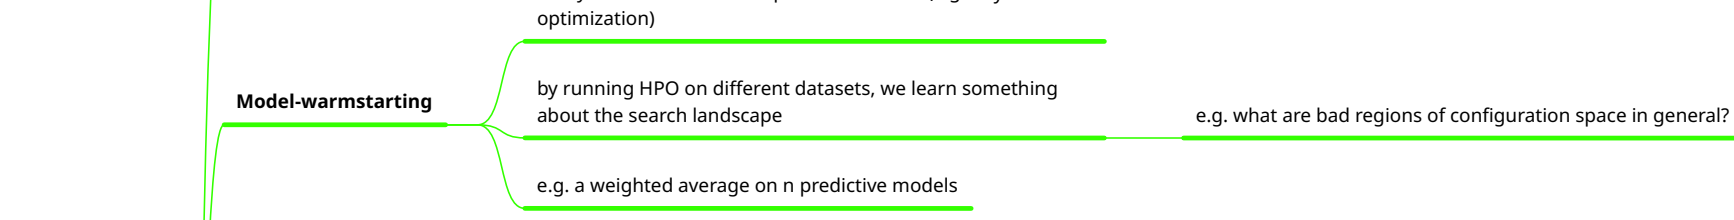


Meta-learning

Warmstarting



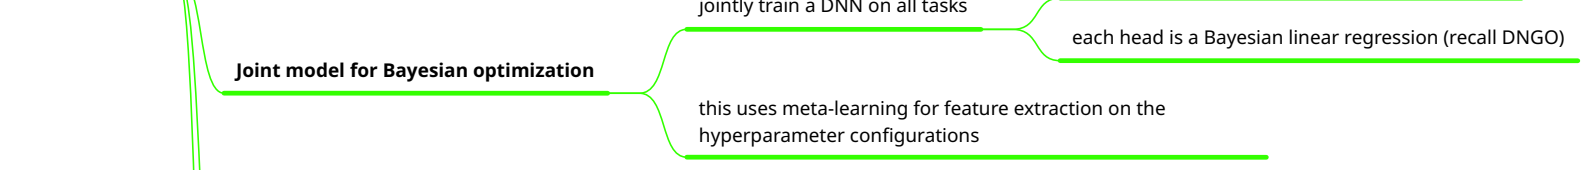
Model warmstarting



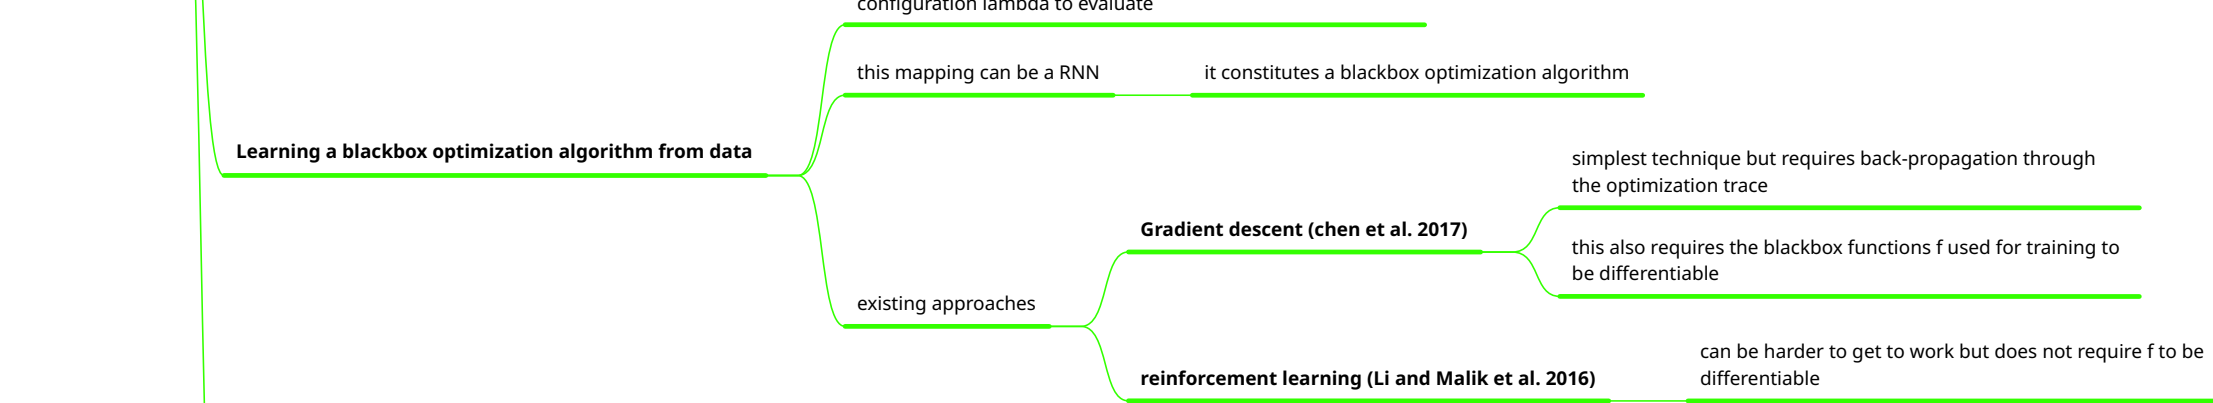
Task-independent Recommendations



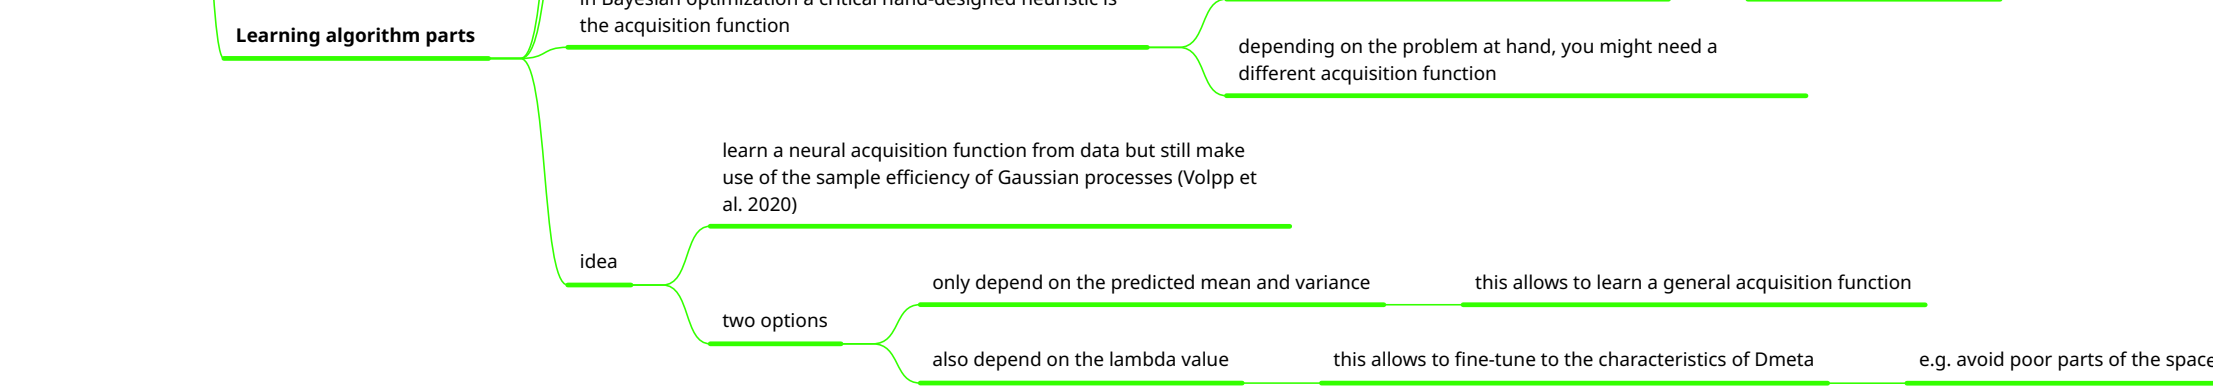
joint model for Bayesian optimization



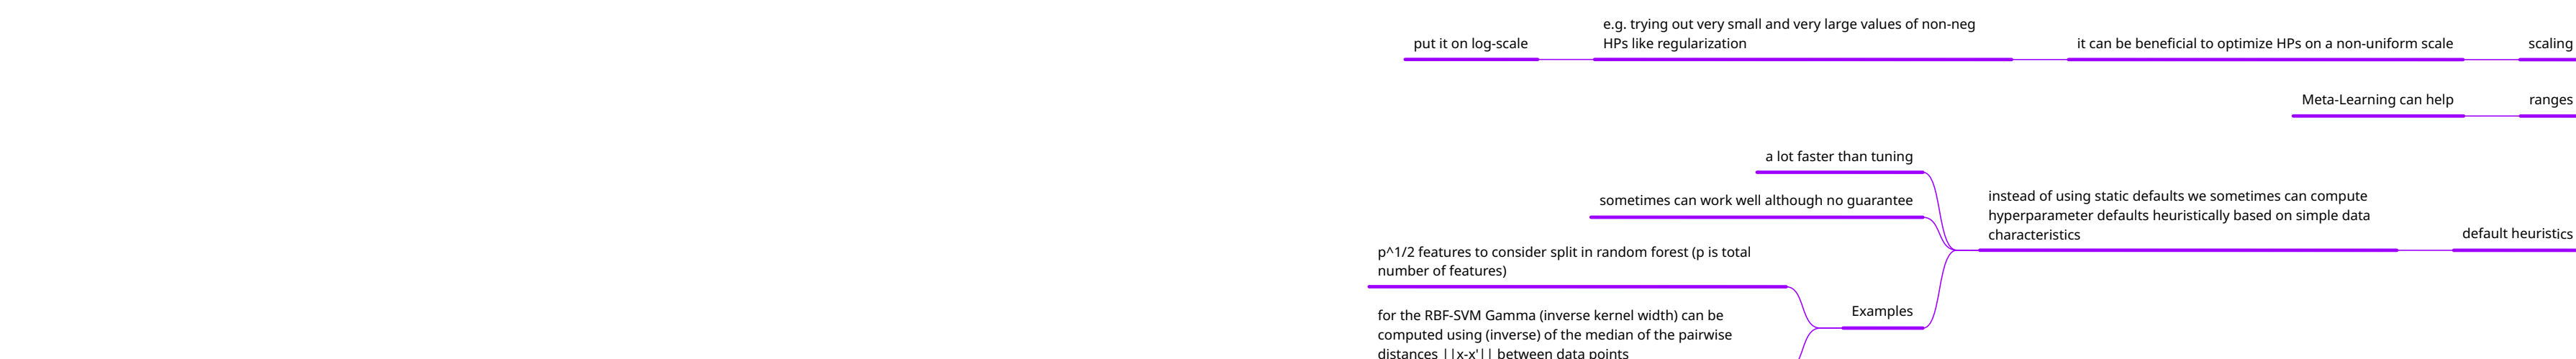
Learning a blackbox optimization algorithm from data



Learning algorithm parts



Tips



Challenges

